

Introduction to ontologies

Semantic similarity

Michel Dumontier & Robert Hoehndorf

ISMB 2018

Overview

1. Ontologies and graphs
2. Semantic similarity
3. Applications
4. Hand-on session

Some examples

- ▶ Are cyclin dependent kinases *functionally* more similar to lipid kinases or to riboflavin kinases? How about *phenotypically*?

Some examples

- ▶ Are cyclin dependent kinases *functionally* more similar to lipid kinases or to riboflavin kinases? How about *phenotypically*?
- ▶ Which protein in the *mouse* is functionally most similar to the zebrafish *gustducin* protein?

Some examples

- ▶ Are cyclin dependent kinases *functionally* more similar to lipid kinases or to riboflavin kinases? How about *phenotypically*?
- ▶ Which protein in the *mouse* is functionally most similar to the zebrafish *gustducin* protein?
- ▶ Which mouse knockout resembles *Bardet-Biedl Syndrome 8*?

Some examples

- ▶ Are cyclin dependent kinases *functionally* more similar to lipid kinases or to riboflavin kinases? How about *phenotypically*?
- ▶ Which protein in the *mouse* is functionally most similar to the zebrafish *gustducin* protein?
- ▶ Which mouse knockout resembles *Bardet-Biedl Syndrome 8*?
- ▶ Are there mouse knockouts that resemble the side effects of diclofenac?

Some examples

- ▶ Are cyclin dependent kinases *functionally* more similar to lipid kinases or to riboflavin kinases? How about *phenotypically*?
- ▶ Which protein in the *mouse* is functionally most similar to the zebrafish *gustducin* protein?
- ▶ Which mouse knockout resembles *Bardet-Biedl Syndrome 8*?
- ▶ Are there mouse knockouts that resemble the side effects of diclofenac?
- ▶ Which genetic disease produces similar symptoms to ebola?

Some examples

- ▶ Are cyclin dependent kinases *functionally* more similar to lipid kinases or to riboflavin kinases? How about *phenotypically*?
- ▶ Which protein in the *mouse* is functionally most similar to the zebrafish *gustducin* protein?
- ▶ Which mouse knockout resembles *Bardet-Biedl Syndrome 8*?
- ▶ Are there mouse knockouts that resemble the side effects of diclofenac?
- ▶ Which genetic disease produces similar symptoms to ebola?
- ▶ Does functional similarity correlate with phenotypic similarity?

Ontologies and graphs

- ▶ semantic similarity measures can be graph-based, feature-based, or model-based
- ▶ we may need to generate graphs from ontologies
 - ▶ *is-a* relations are easy
 - ▶ how about *part-of*, *regulates*, *precedes*, etc.?
- ▶ relational patterns are defined in OBO Relation Ontology
 - ▶ in first order logic
 - ▶ needs to translate them into OWL

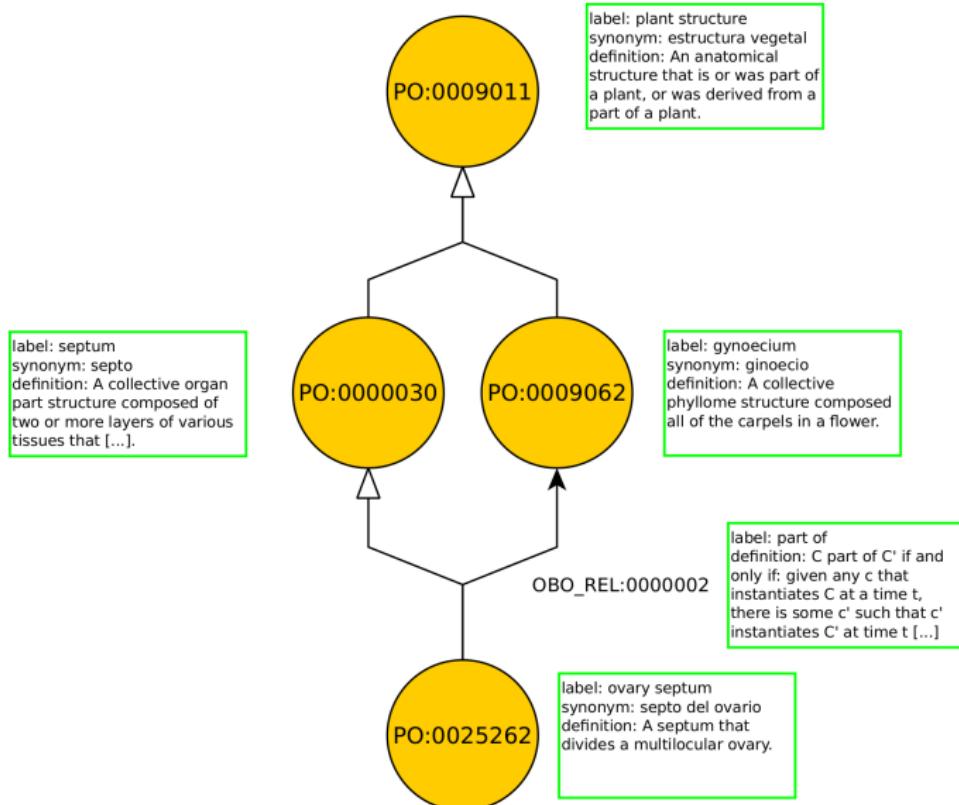
Relations as patterns

- ▶ X SubClassOf: Y: $X \xrightarrow{\text{is-a}} Y$
- ▶ X SubClassOf: part-of some Y: $X \xrightarrow{\text{part-of}} Y$
- ▶ X SubClassOf: regulates some Y: $X \xrightarrow{\text{regulates}} Y$
- ▶ X DisjointWith: Y: $X \xleftarrow{\text{disjoint}} Y$
- ▶ X EquivalentTo: Y: $X \xrightleftharpoons{\equiv} Y, \{X, Y\}$

Relations as patterns

- ▶ OBO Relation Ontology (RO):
 - ▶ <https://github.com/oborel/obo-relations>
- ▶ Basic Formal Ontology (BFO):
 - ▶ provides top-level classes
 - ▶ Continuant, Process, Function, Material object, etc.
 - ▶ used for some OBO Foundry ontologies
- ▶ RO and BFO provide a top-level system of classes and relations shared across many biomedical ontologies
 - ▶ even GO, although somewhat hidden!

Relations as patterns



Relations as patterns: Quiz

- ▶ *classes*

- a represent kinds of things in the world
- b represent ideas in people's heads
- c represent words

Relations as patterns: Quiz

- ▶ *classes*
 - a represent kinds of things in the world
 - b represent ideas in people's heads
 - c represent words
- ▶ *instances* of a class are individuals that
 - a satisfy the class intension
 - b satisfy the axioms used to specify the class
 - c are always concrete, material entities

Relations as patterns: Quiz

- ▶ *classes*
 - a represent kinds of things in the world
 - b represent ideas in people's heads
 - c represent words
- ▶ *instances* of a class are individuals that
 - a satisfy the class intension
 - b satisfy the axioms used to specify the class
 - c are always concrete, material entities
- ▶ *relations between classes* are
 - a interactions between instances of the class
 - b abbreviations of axioms that constrain two classes
 - c relations between ideas people have about the classes

Relations as patterns: Quiz

- ▶ *classes*
 - a represent kinds of things in the world
 - b represent ideas in people's heads
 - c represent words
- ▶ *instances* of a class are individuals that
 - a satisfy the class intension
 - b satisfy the axioms used to specify the class
 - c are always concrete, material entities
- ▶ *relations between classes* are
 - a interactions between instances of the class
 - b abbreviations of axioms that constrain two classes
 - c relations between ideas people have about the classes
- ▶ *axioms* are
 - a specification of conditions that instances of classes must satisfy
 - b rules that can be executed to produce new knowledge
 - c statements that are considered to be true in a domain of knowledge

How to measure similarity?

- ▶ semantic similarity measures similarity between classes
- ▶ semantic similarity measures similarity between instances of classes
- ▶ semantic similarity measures similarity between entities annotated with classes
- ▶ ⇒ reduce all of this to similarity between classes

How to measure similarity?

What properties do we want in a similarity measure?

A function $\text{sim} : D \times D$ is a similarity on D if, for all $x, y \in D$, the function sim is:

How to measure similarity?

What properties do we want in a similarity measure?

A function $\text{sim} : D \times D$ is a similarity on D if, for all $x, y \in D$, the function sim is:

- ▶ non-negative: $\text{sim}(x, y) \geq 0$ for all x, y

How to measure similarity?

What properties do we want in a similarity measure?

A function $\text{sim} : D \times D$ is a similarity on D if, for all $x, y \in D$, the function sim is:

- ▶ non-negative: $\text{sim}(x, y) \geq 0$ for all x, y
- ▶ symmetric: $\text{sim}(x, y) = \text{sim}(y, x)$

How to measure similarity?

What properties do we want in a similarity measure?

A function $\text{sim} : D \times D$ is a similarity on D if, for all $x, y \in D$, the function sim is:

- ▶ non-negative: $\text{sim}(x, y) \geq 0$ for all x, y
- ▶ symmetric: $\text{sim}(x, y) = \text{sim}(y, x)$
- ▶ reflexive: $\text{sim}(x, x) = \max_D$

How to measure similarity?

What properties do we want in a similarity measure?

A function $\text{sim} : D \times D$ is a similarity on D if, for all $x, y \in D$, the function sim is:

- ▶ non-negative: $\text{sim}(x, y) \geq 0$ for all x, y
- ▶ symmetric: $\text{sim}(x, y) = \text{sim}(y, x)$
- ▶ reflexive: $\text{sim}(x, x) = \max_D$
 - ▶ weaker form: $\text{sim}(x, x) > \text{sim}(x, y)$ for all $x \neq y$

How to measure similarity?

What properties do we want in a similarity measure?

A function $\text{sim} : D \times D$ is a similarity on D if, for all $x, y \in D$, the function sim is:

- ▶ non-negative: $\text{sim}(x, y) \geq 0$ for all x, y
- ▶ symmetric: $\text{sim}(x, y) = \text{sim}(y, x)$
- ▶ reflexive: $\text{sim}(x, x) = \max_D$
 - ▶ weaker form: $\text{sim}(x, x) > \text{sim}(x, y)$ for all $x \neq y$
- ▶ $\text{sim}(x, x) > \text{sim}(x, y)$ for $x \neq y$

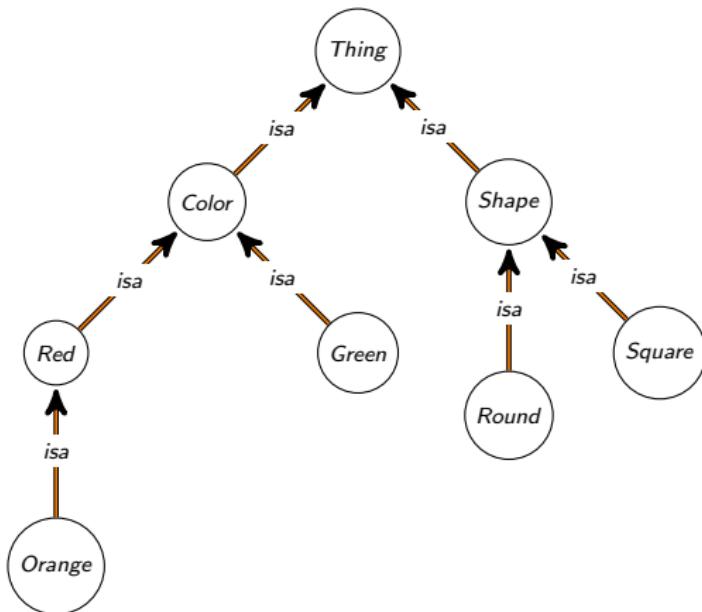
How to measure similarity?

What properties do we want in a similarity measure?

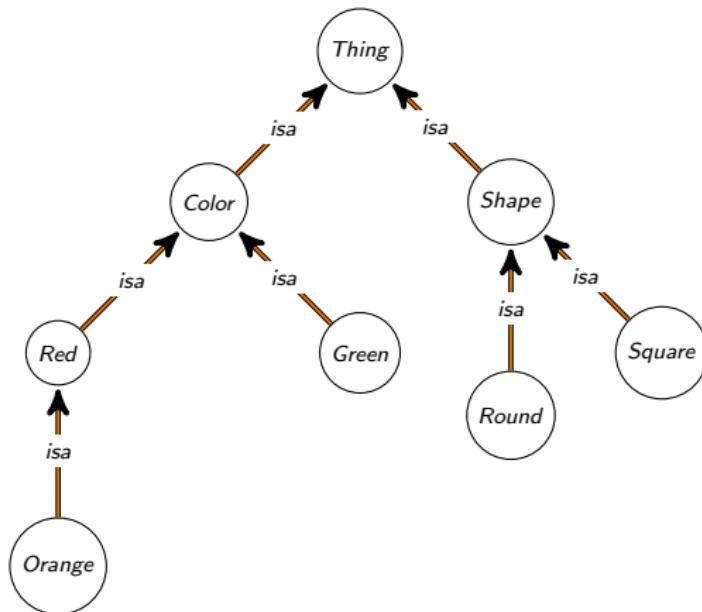
A function $\text{sim} : D \times D$ is a similarity on D if, for all $x, y \in D$, the function sim is:

- ▶ non-negative: $\text{sim}(x, y) \geq 0$ for all x, y
- ▶ symmetric: $\text{sim}(x, y) = \text{sim}(y, x)$
- ▶ reflexive: $\text{sim}(x, x) = \max_D$
 - ▶ weaker form: $\text{sim}(x, x) > \text{sim}(x, y)$ for all $x \neq y$
- ▶ $\text{sim}(x, x) > \text{sim}(x, y)$ for $x \neq y$
- ▶ sim is a *normalized* similarity measure if it has values in $[0, 1]$

How to measure similarity?

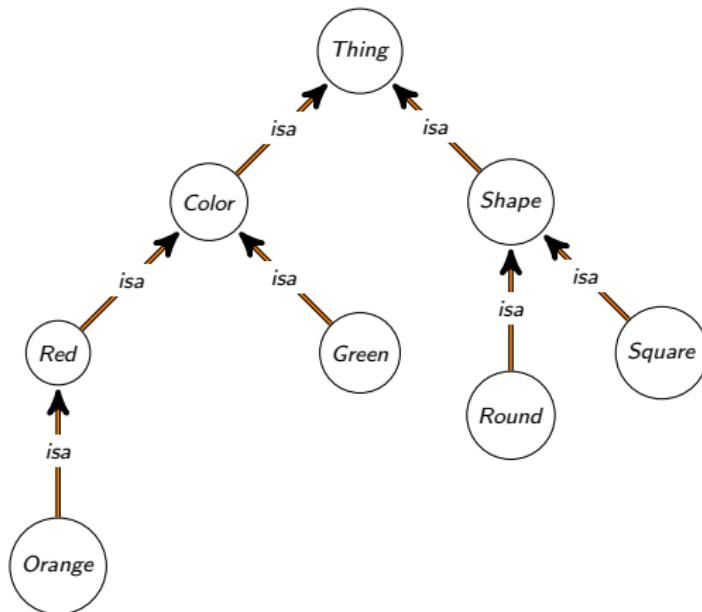


How to measure similarity?



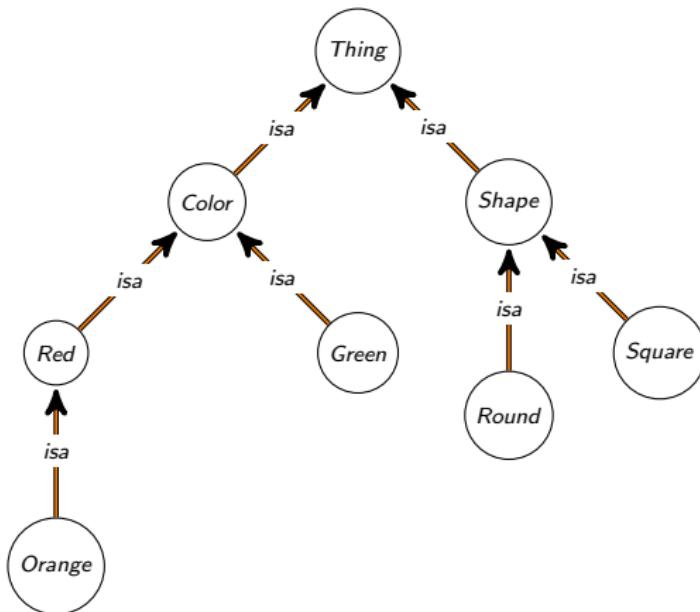
- ▶ distance on shortest path (Rada *et al.*, 1989)

How to measure similarity?



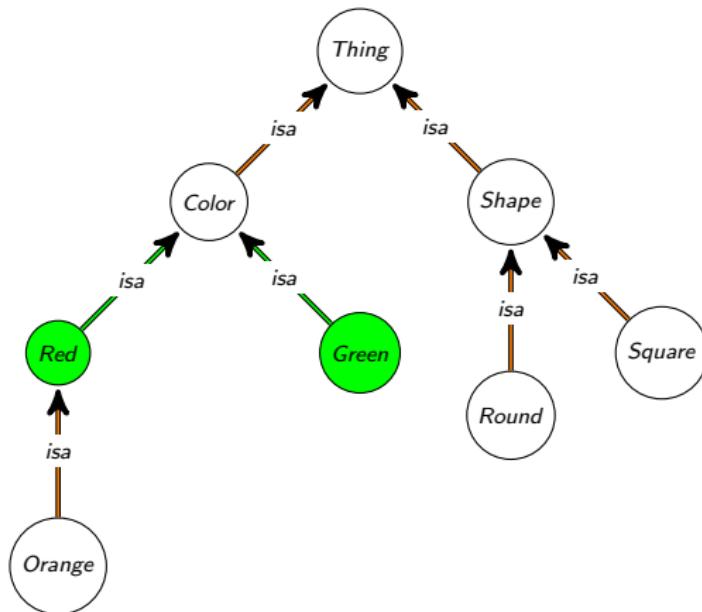
- ▶ distance on shortest path (Rada *et al.*, 1989)
- ▶ $dist_{Rada}(u, v) = sp(u, \text{isa}, v)$

How to measure similarity?



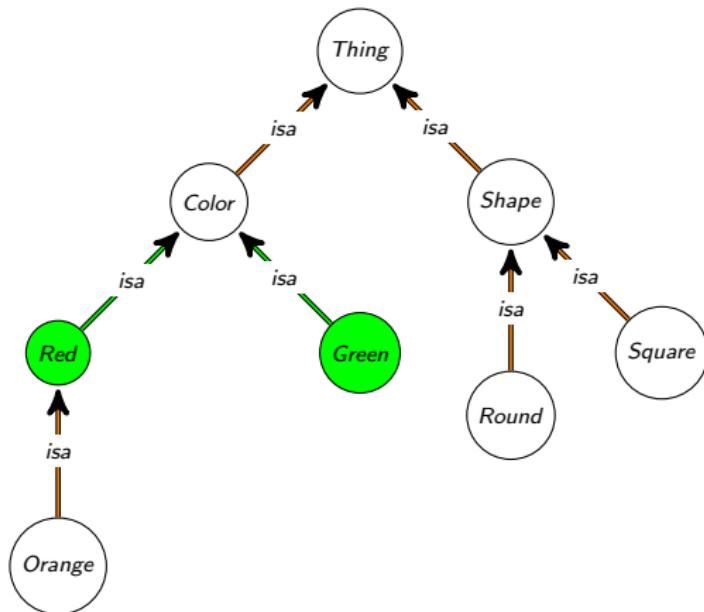
- ▶ distance on shortest path (Rada *et al.*, 1989)
- ▶ $dist_{Rada}(u, v) = sp(u, \text{isa}, v)$
- ▶ $sim_{Rada}(u, v) = \frac{1}{dist_{Rada}(u, v) + 1}$

How to measure similarity?



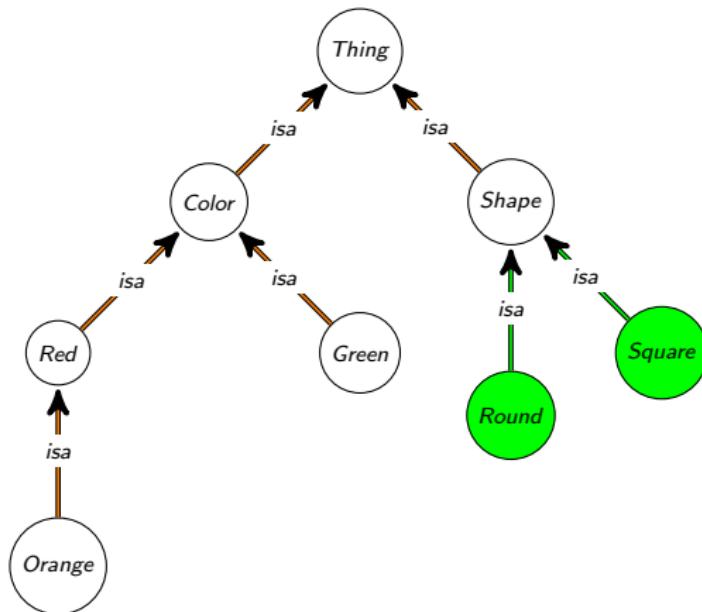
- ▶ distance on shortest path

How to measure similarity?



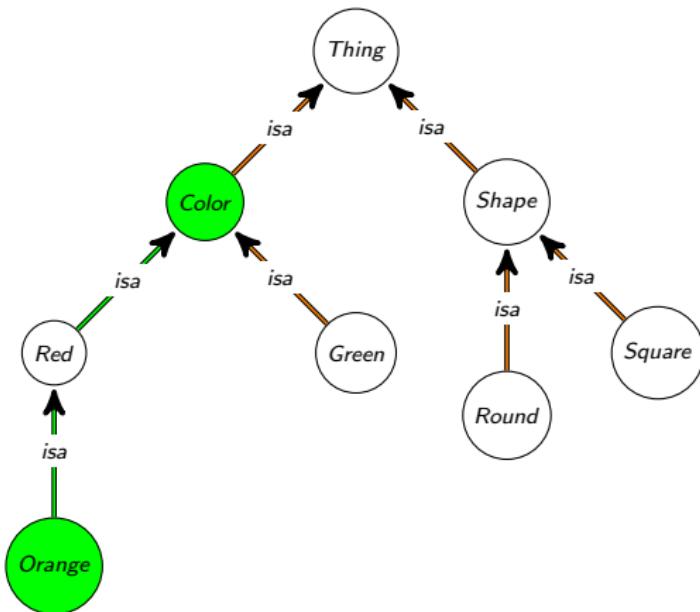
- distance on shortest path
- distance(green, red) = 2
- $sim_{Rada}(green, red) = \frac{1}{3}$

How to measure similarity?



- ▶ distance on shortest path
- ▶ $\text{distance}(\text{square}, \text{round}) = 2$
- ▶ $\text{sim}_{\text{Rada}}(\text{square}, \text{round}) = \frac{1}{3}$

How to measure similarity?



- ▶ distance on shortest path
- ▶ $\text{distance}(\text{orange}, \text{color}) = 2$
- ▶ $\text{sim}_{\text{Rada}}(\text{orange}, \text{color}) = \frac{1}{3}$

How to measure similarity?

- ▶ shortest path is not always intuitive

How to measure similarity?

- ▶ shortest path is not always intuitive
- ▶ we need a way to determine *specificity* of a class
 - ▶ number of ancestors
 - ▶ number of children
 - ▶ information content

How to measure similarity?

- ▶ shortest path is not always intuitive
- ▶ we need a way to determine *specificity* of a class
 - ▶ number of ancestors
 - ▶ number of children
 - ▶ information content
- ▶ *density* of a branch in the ontology
 - ▶ number of siblings
 - ▶ information content

How to measure similarity?

- ▶ shortest path is not always intuitive
- ▶ we need a way to determine *specificity* of a class
 - ▶ number of ancestors
 - ▶ number of children
 - ▶ information content
- ▶ *density* of a branch in the ontology
 - ▶ number of siblings
 - ▶ information content
- ▶ account for different edge types
 - ▶ non-uniform edge weighting

How to measure similarity

- ▶ term specificity measure $\sigma : C \mapsto \mathbb{R}$:
 - ▶ $x \sqsubseteq y \rightarrow \sigma(x) \geq \sigma(y)$

How to measure similarity

- ▶ term specificity measure $\sigma : C \mapsto \mathbb{R}$:

- ▶ $x \sqsubseteq y \rightarrow \sigma(x) \geq \sigma(y)$

- ▶ intrinsic:

- ▶ $\sigma(x) = f(\text{depth}(x))$

- ▶ $\sigma(x) = f(A(x))$ (for ancestors $A(x)$)

- ▶ $\sigma(x) = f(D(x))$ (for descendants $D(x)$)

- ▶ many more, e.g., Zhou et al.:

$$\sigma(x) = k \cdot \left(1 - \frac{\log |D(x)|}{\log |C|}\right) + (1 - k) \frac{\log \text{depth}(x)}{\log \text{depth}(G_T)}$$

How to measure similarity

- ▶ term specificity measure $\sigma : C \mapsto \mathbb{R}$:

- ▶ $x \sqsubseteq y \rightarrow \sigma(x) \geq \sigma(y)$

- ▶ intrinsic:

- ▶ $\sigma(x) = f(\text{depth}(x))$

- ▶ $\sigma(x) = f(A(x))$ (for ancestors $A(x)$)

- ▶ $\sigma(x) = f(D(x))$ (for descendants $D(x)$)

- ▶ many more, e.g., Zhou et al.:

$$\sigma(x) = k \cdot \left(1 - \frac{\log |D(x)|}{\log |C|}\right) + (1 - k) \frac{\log \text{depth}(x)}{\log \text{depth}(G_T)}$$

- ▶ extrinsic:

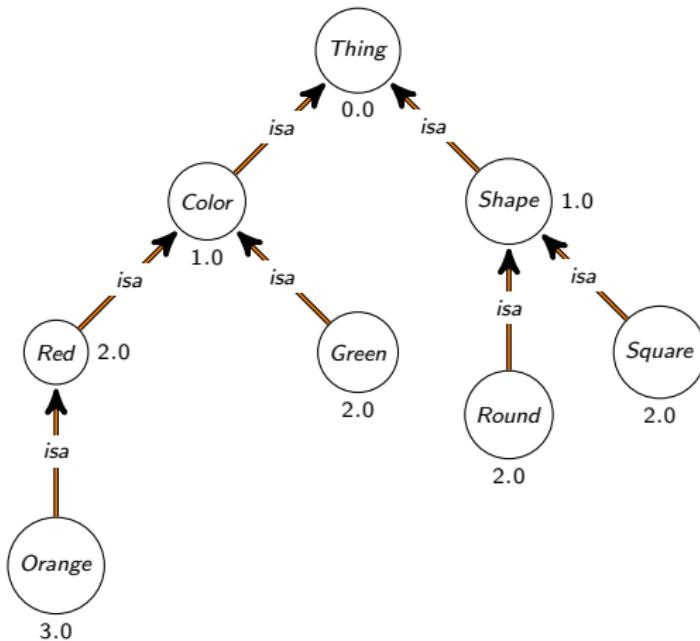
- ▶ $\sigma(x)$ defined as a function of instances (or annotations) I

- ▶ note: the number of instances monotonically decreases with increasing depth in taxonomies

- ▶ Resnik 1995: $eIC_{Resnik}(x) = -\log p(x)$ (with $p(x) = \frac{|I(x)|}{|I|}$)

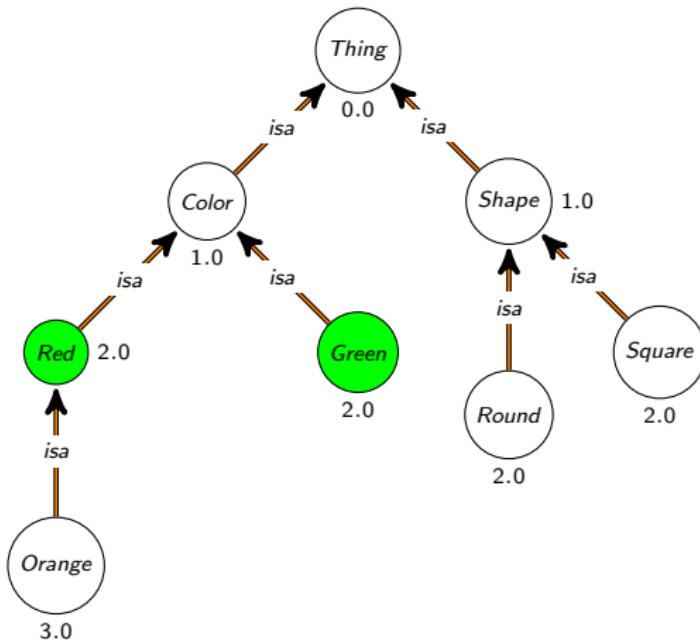
- ▶ in biology, one of the most popular specificity measure when annotations are present

How to measure similarity?



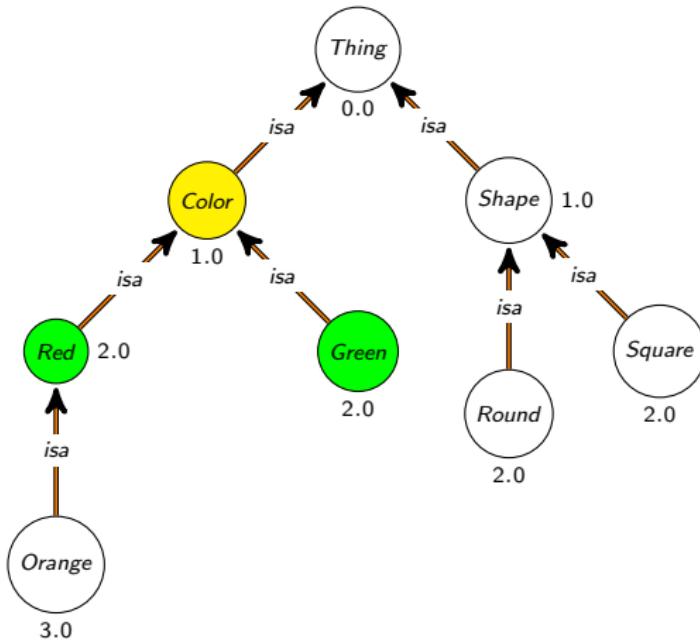
- ▶ Resnik 1995:
similarity between x and y is the
information content
of the *most
informative common
ancestor*

How to measure similarity?



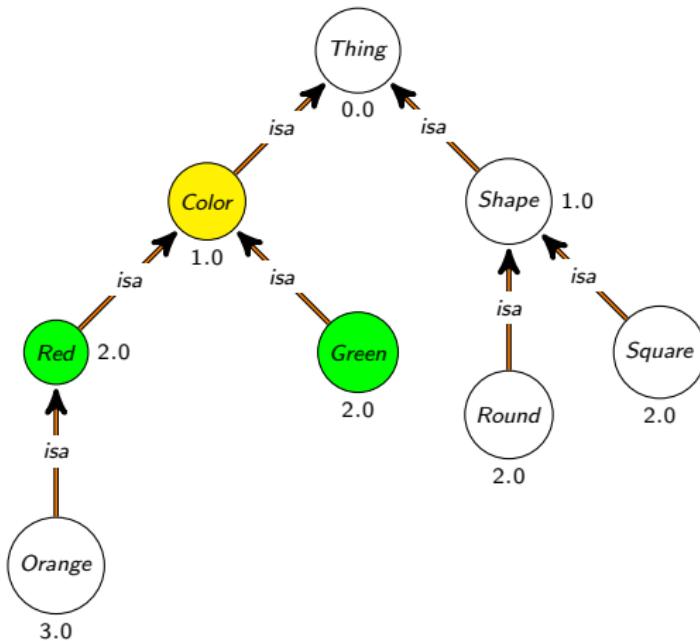
- ▶ Resnik 1995:
similarity between x and y is the
information content
of the *most
informative common
ancestor*

How to measure similarity?



- ▶ Resnik 1995:
similarity between x and y is the
information content
of the *most
informative common
ancestor*

How to measure similarity?

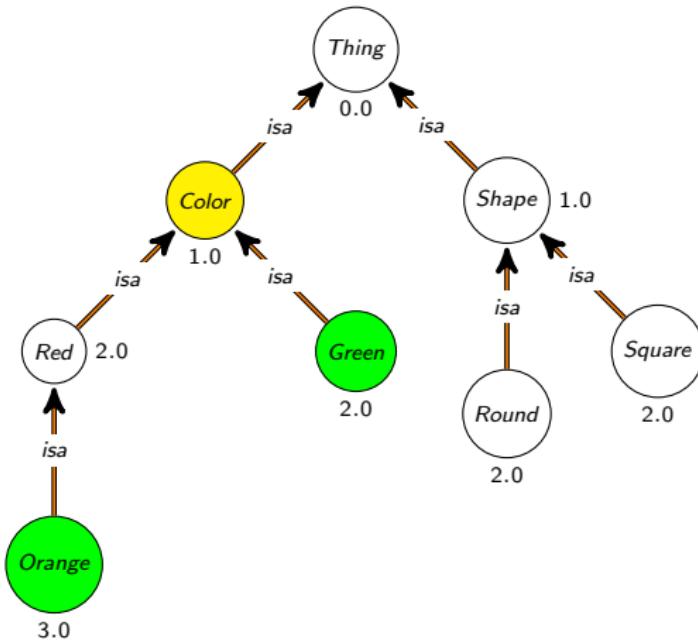


► Resnik 1995:
similarity between x and y is the
information content
of the *most
informative common
ancestor*



$$sim_{Resnik}(Green, Red) = 1.0$$

How to measure similarity?

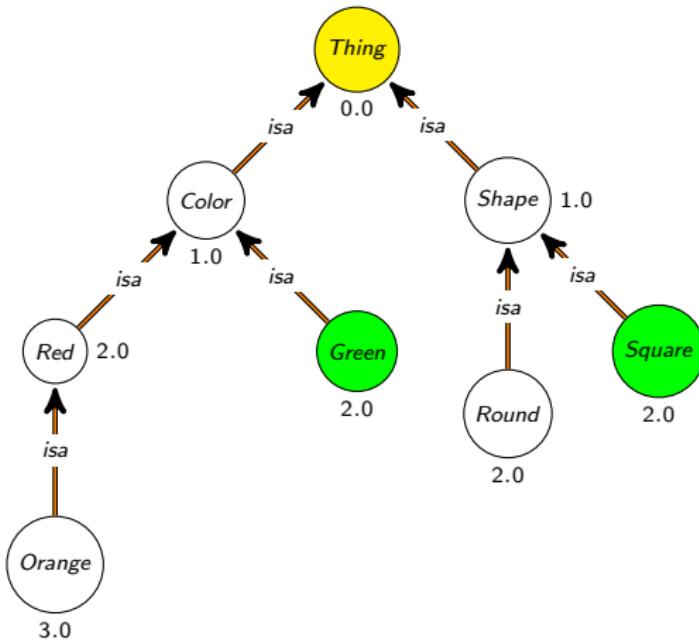


► Resnik 1995:
similarity between x and y is the
information content
of the *most
informative common
ancestor*



$$sim_{Resnik}(Green, Orange) = 1.0$$

How to measure similarity?



► Resnik 1995:
similarity between x and y is the
information content
of the *most
informative common
ancestor*

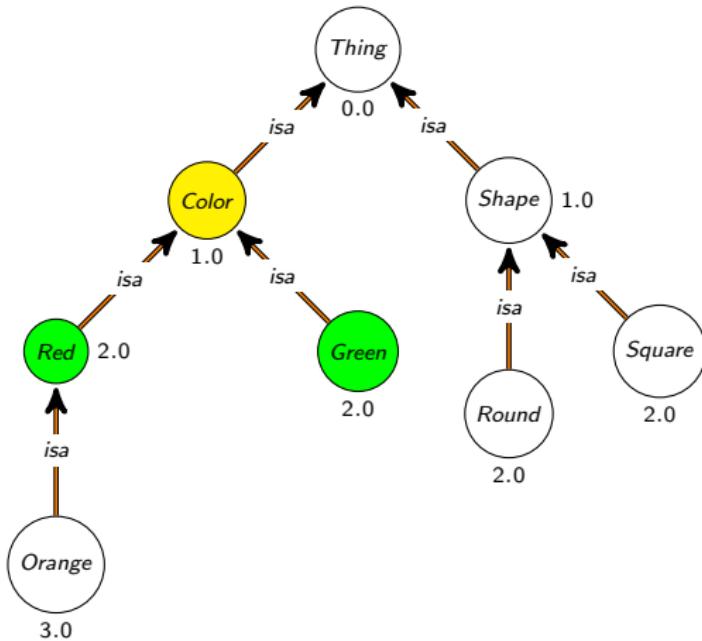


$$sim_{Resnik}(Square, Orange) \\ 0.0$$

How to measure similarity?

- ▶ (Red, Green) and (Orange, Green) have the same similarity
- ▶ need to incorporate the specificity of the compared classes

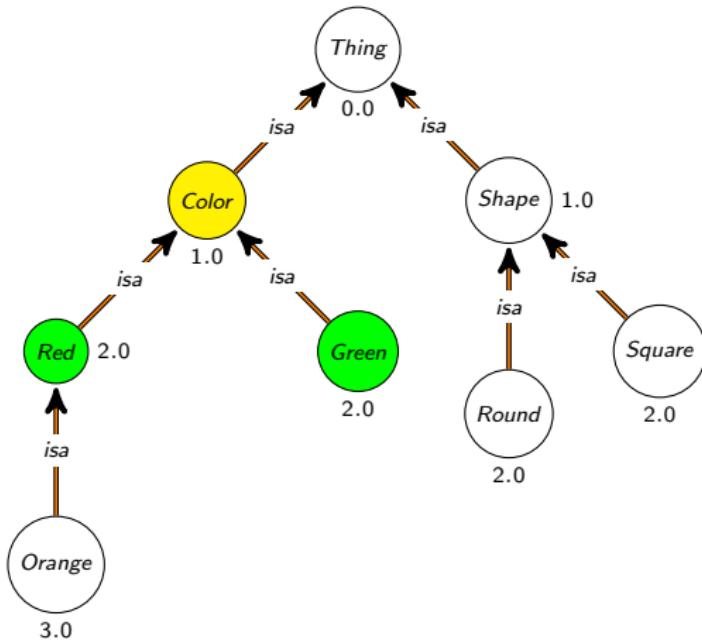
How to measure similarity?



► Lin 1998:

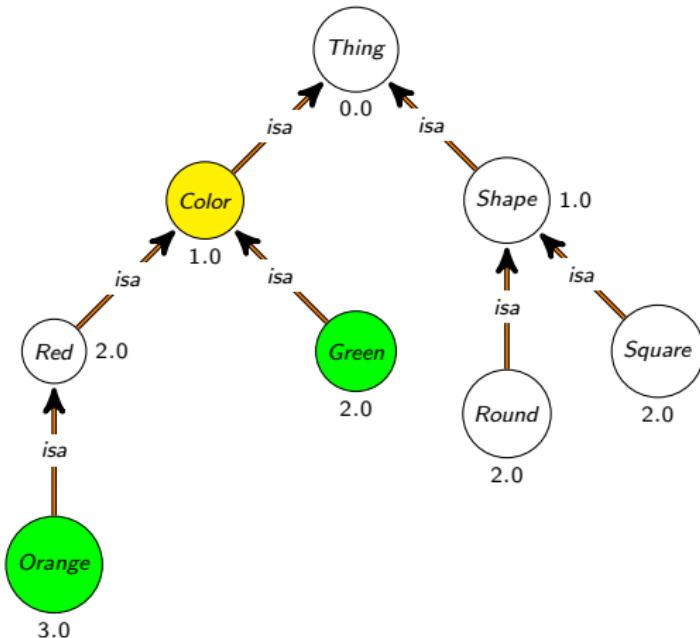
$$sim_{Lin}(x, y) = \frac{2 \cdot IC(MICA(x, y))}{IC(x) + IC(y)}$$

How to measure similarity?



- Lin 1998:
$$sim_{Lin}(x, y) = \frac{2 \cdot IC(MICA(x, y))}{IC(x) + IC(y)}$$
- $sim_{Lin}(Green, Red) = 0.5$

How to measure similarity?



► Lin 1998:

$$sim_{Lin}(x, y) = \frac{2 \cdot IC(MICA(x, y))}{IC(x) + IC(y)}$$

►

$$sim_{Lin}(Green, Orange) = 0.4$$

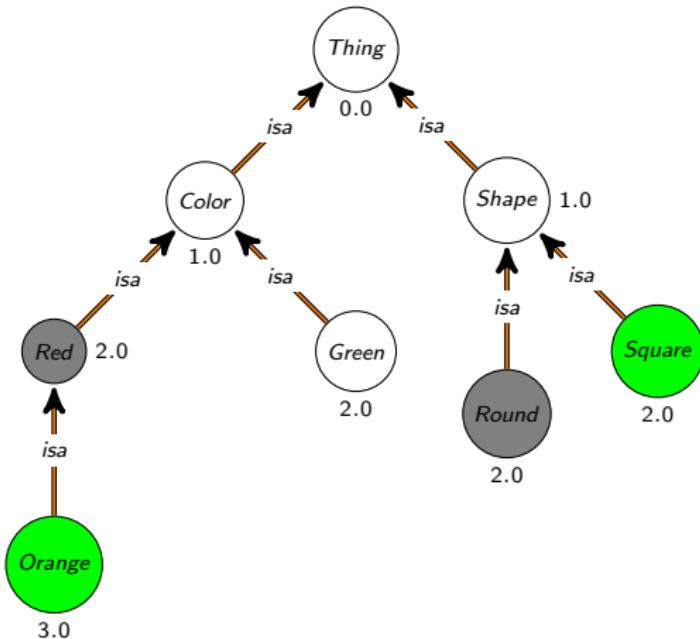
How to measure similarity?

- ▶ many(!) others:
 - ▶ Jiang & Conrath 1997
 - ▶ Mazandu & Mulder 2013
 - ▶ Schlicker et al. 2009
 - ▶ ...

How to measure similarity?

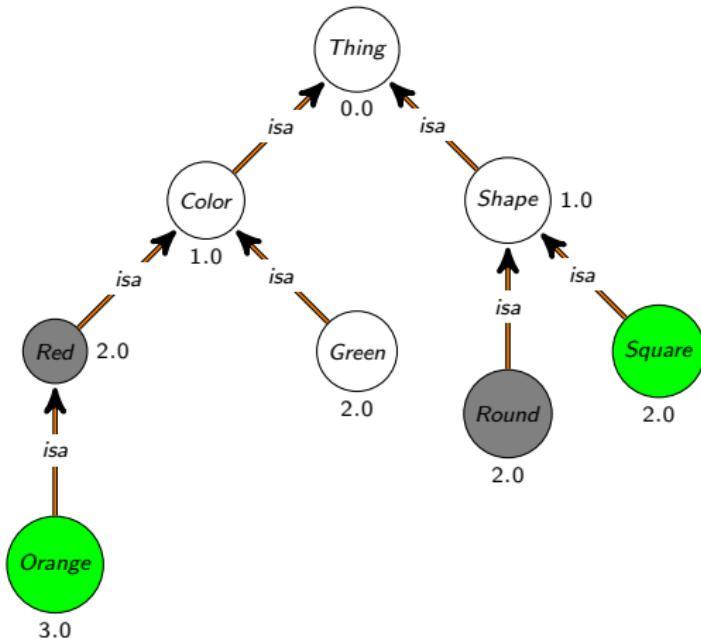
- ▶ we only looked at comparing pairs of classes
- ▶ mostly, we want to compare *sets* of classes
 - ▶ set of GO annotations
 - ▶ set of signs and symptoms
 - ▶ set of phenotypes
- ▶ two approaches:
 - ▶ compare each class individually, then merge
 - ▶ directly set-based similarity measures

How to measure similarity?



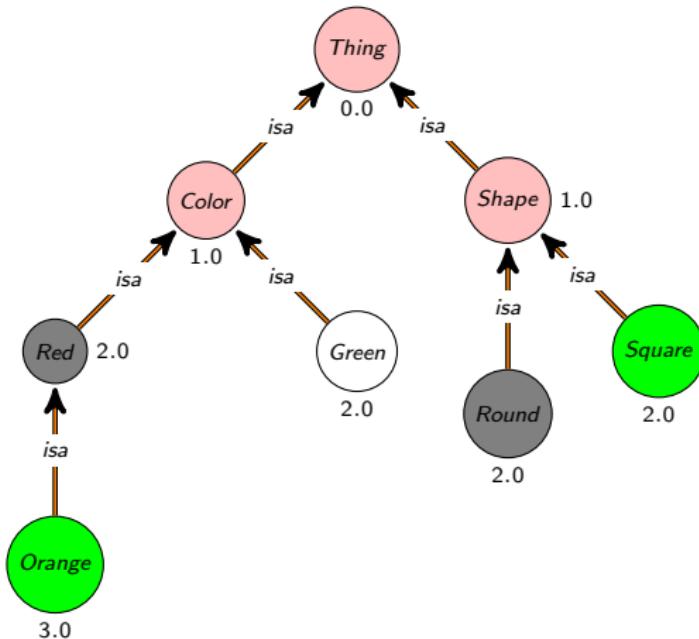
- ▶ similarity between a square-and-orange thing and a round-and-red thing

How to measure similarity?



- ▶ similarity between a square-and-orange thing and a round-and-red thing
- ▶ Pesquita et al., 2007:
$$\text{simGIC}(X, Y) = \frac{\sum_{c \in A(X) \cap A(Y)} IC(c)}{\sum_{c \in A(X) \cup A(Y)} IC(c)}$$

How to measure similarity?



- ▶ similarity between a square-and-orange thing and a round-and-red thing
- ▶ Pesquita et al., 2007:
$$\text{simGIC}(X, Y) = \frac{\sum_{c \in A(X) \cap A(Y)} IC(c)}{\sum_{c \in A(X) \cup A(Y)} IC(c)}$$
- ▶ $\text{simGIC}(so, rr) = \frac{2}{11}$

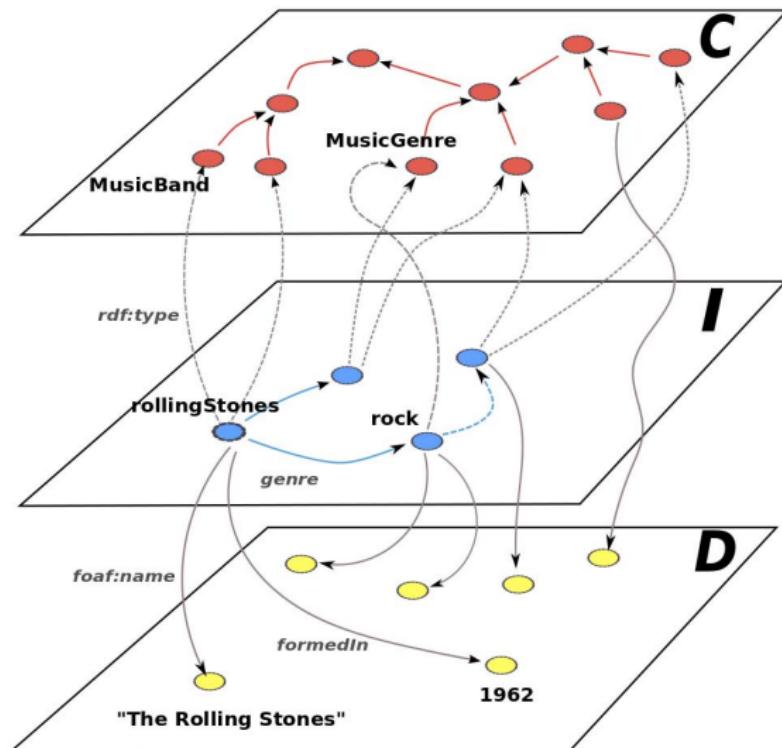
How to measure similarity?

- ▶ alternatively: use different merging strategies
- ▶ common: average, maximum, **best-matching average**
 - ▶ Average: $sim_A(X, Y) = \frac{\sum_{x \in X} \sum_{y \in Y} sim(x, y)}{|X| \times |Y|}$
 - ▶ Max average: $sim_{MA}(X, Y) = \frac{1}{|X|} \sum_{x \in X} \max_{y \in Y} sim(x, y)$
 - ▶ Best match average: $sim_{BMA}(X, Y) = \frac{sim_{MA}(X, Y) + sim_{MA}(Y, X)}{2}$

How to measure similarity?

- ▶ Semantic Measures Library:
 - ▶ comprehensive Java library
 - ▶ <http://www.semantic-measures-library.org/>
- ▶ R packages: GOSim, GOSemSim, HPOSim, LSAfun, ontologySimilarity,...
- ▶ Python: sematch, fastsemsim (GO only)

How to measure similarity?



From Harispe et al., Semantic Similarity From Natural Language And Ontology Analysis, 2015.

How to measure similarity?

- ▶ Shortest Path
 - ▶ applicable to arbitrary knowledge graphs
 - ▶ does not capture similarity well over all edge types, e.g., *disjointWith*, *differentFrom*, *opposite-of*, etc.
- ▶ Random Walk
 - ▶ with or without restart
 - ▶ iterated
 - ▶ does not consider edge labels ⇒ captures only adjacency of nodes
 - ▶ scores whole graph with *probability* of being in a state
 - ▶ can take multiple seed nodes
 - ▶ widely used to find disease genes

How to measure similarity?

- ▶ feature learning on knowledge graph

How to measure similarity?

- ▶ feature learning on knowledge graph
- ▶ e.g., iterated, edge-labeled random walk
 - ▶ over instances and classes
 - ▶ walks form *sentences*
 - ▶ sentences form a *corpus*
 - ▶ feature learning on corpus through Word2Vec (or factorization of co-occurrence matrix)
 - ▶ RDF2Vec: <http://data.dws.informatik.uni-mannheim.de/rdf2vec/>

How to measure similarity?

- ▶ feature learning on knowledge graph
- ▶ e.g., iterated, edge-labeled random walk
 - ▶ over instances and classes
 - ▶ walks form *sentences*
 - ▶ sentences form a *corpus*
 - ▶ feature learning on corpus through Word2Vec (or factorization of co-occurrence matrix)
 - ▶ RDF2Vec: <http://data.dws.informatik.uni-mannheim.de/rdf2vec/>
 - ▶ with support for reasoning over bio-ontologies:
<https://github.com/bio-ontology-research-group/walking-rdf-and-owl>
 - ▶ (self-promotion warning)

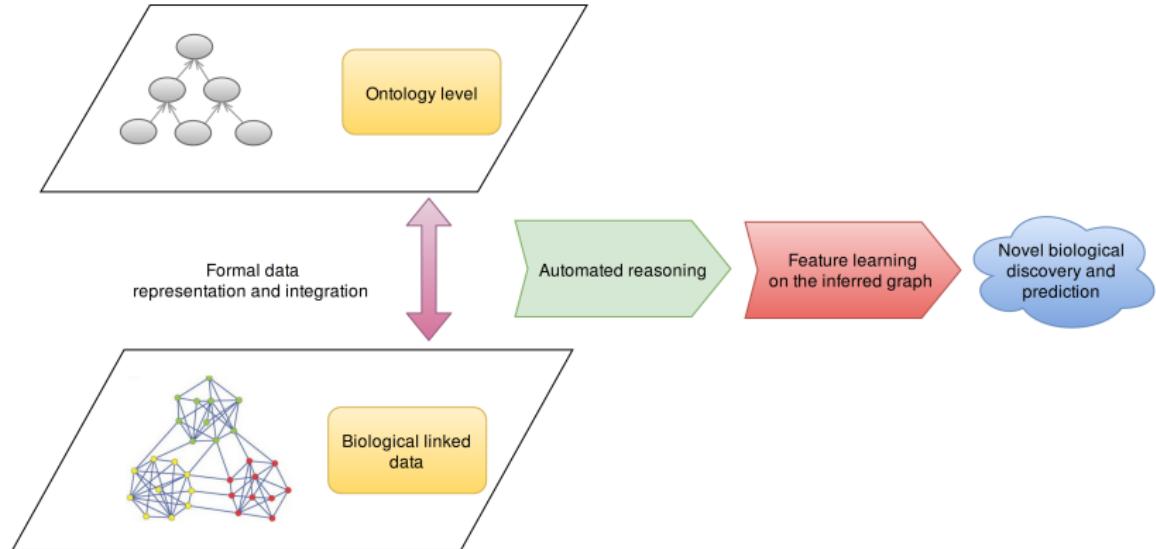
How to measure similarity?

- ▶ feature learning on knowledge graph
- ▶ e.g., iterated, edge-labeled random walk
 - ▶ over instances and classes
 - ▶ walks form *sentences*
 - ▶ sentences form a *corpus*
 - ▶ feature learning on corpus through Word2Vec (or factorization of co-occurrence matrix)
 - ▶ RDF2Vec: <http://data.dws.informatik.uni-mannheim.de/rdf2vec/>
 - ▶ with support for reasoning over bio-ontologies:
<https://github.com/bio-ontology-research-group/walking-rdf-and-owl>
 - ▶ (self-promotion warning)
- ▶ Translational knowledge graph embeddings: TransE, TransR, TransE, HolE, etc.
 - ▶ analogy-based
 - ▶ <https://github.com/thunlp/KB2E>

How to measure similarity?

- ▶ feature learning on knowledge graph
- ▶ e.g., iterated, edge-labeled random walk
 - ▶ over instances and classes
 - ▶ walks form *sentences*
 - ▶ sentences form a *corpus*
 - ▶ feature learning on corpus through Word2Vec (or factorization of co-occurrence matrix)
 - ▶ RDF2Vec: <http://data.dws.informatik.uni-mannheim.de/rdf2vec/>
 - ▶ with support for reasoning over bio-ontologies:
<https://github.com/bio-ontology-research-group/walking-rdf-and-owl>
 - ▶ (self-promotion warning)
- ▶ Translational knowledge graph embeddings: TransE, TransR, TransE, HolE, etc.
 - ▶ analogy-based
 - ▶ <https://github.com/thunlp/KB2E>
- ▶ generates (dense) feature vectors for nodes (classes, instances) and relations

How to measure similarity?



How to measure similarity?

- ▶ vector-based similarity measure
- ▶ cosine similarity: $sim(X, Y) = \frac{\sum_{i=1}^n X_i Y_i}{\sqrt{\sum_{i=1}^n X_i^2} \sqrt{\sum_{i=1}^n Y_i^2}}$
 - ▶ bounded between $[-1, 1]$
- ▶ Euclidean distance: $sim(X, Y) = \sqrt{\sum_{i=1}^n (X_i - Y_i)^2}$
 - ▶ not bounded (and rarely used)

How to measure similarity?

- ▶ many graph based semantic similarity measures for comparing two classes
- ▶ several set-based measures
 - ▶ directly set-based
 - ▶ merging pair-wise comparison
- ▶ most useful when comparing instances/annotations
- ▶ other approaches consider relations between instances:
 - ▶ path-based
 - ▶ random-walk
- ▶ very recent: knowledge graph embeddings
 - ▶ and any vector-based similarity measure

How to measure similarity?

Recommended reading:

- ▶ recommended, comprehensive overview: Sebastian Harispe et al. Semantic Similarity from Natural Language and Ontology Analysis. Morgan & Claypool Publishers, 2015
- ▶ Catia Pesquita et al. Semantic Similarity in Biomedical Ontologies. PLoS CB, 2009.
- ▶ Maximilian Nickel et al. A Review of Relational Machine Learning for Knowledge Graphs, Proceedings of the IEEE, 2016.

How to measure similarity: Quiz

- ▶ How many semantic similarity measures are there?
 - a One (and it is called The Semantic Similarity Measure)
 - b Three (graph-based, set-based, feature-based)
 - c Many (depending on context, many functions can determine similarity)

How to measure similarity: Quiz

- ▶ How many semantic similarity measures are there?
 - a One (and it is called The Semantic Similarity Measure)
 - b Three (graph-based, set-based, feature-based)
 - c Many (depending on context, many functions can determine similarity)
- ▶ Specificity of an ontology class
 - a depends on the number of children and ancestors, and the depth
 - b depends on the number of instances (or annotations)
 - c can improve similarity estimates significantly

How to measure similarity: Quiz

- ▶ How many semantic similarity measures are there?
 - a One (and it is called The Semantic Similarity Measure)
 - b Three (graph-based, set-based, feature-based)
 - c Many (depending on context, many functions can determine similarity)
- ▶ Specificity of an ontology class
 - a depends on the number of children and ancestors, and the depth
 - b depends on the number of instances (or annotations)
 - c can improve similarity estimates significantly
- ▶ In the presence of (relations between) instances, semantic similarity
 - a cannot be computed, it only works with ontologies
 - b can be estimated using only class specificity measures
 - c can be computed using feature learning and knowledge graph embeddings

Applications of semantic similarity

- ▶ ontologies are used *almost everywhere* in biology
- ▶ many applications of semantic similarity:
 - ▶ predicting interacting proteins
 - ▶ predict candidate genes
 - ▶ using the guilt-by-association principle, or without
 - ▶ predict drug targets and indications
 - ▶ as features in machine learning models

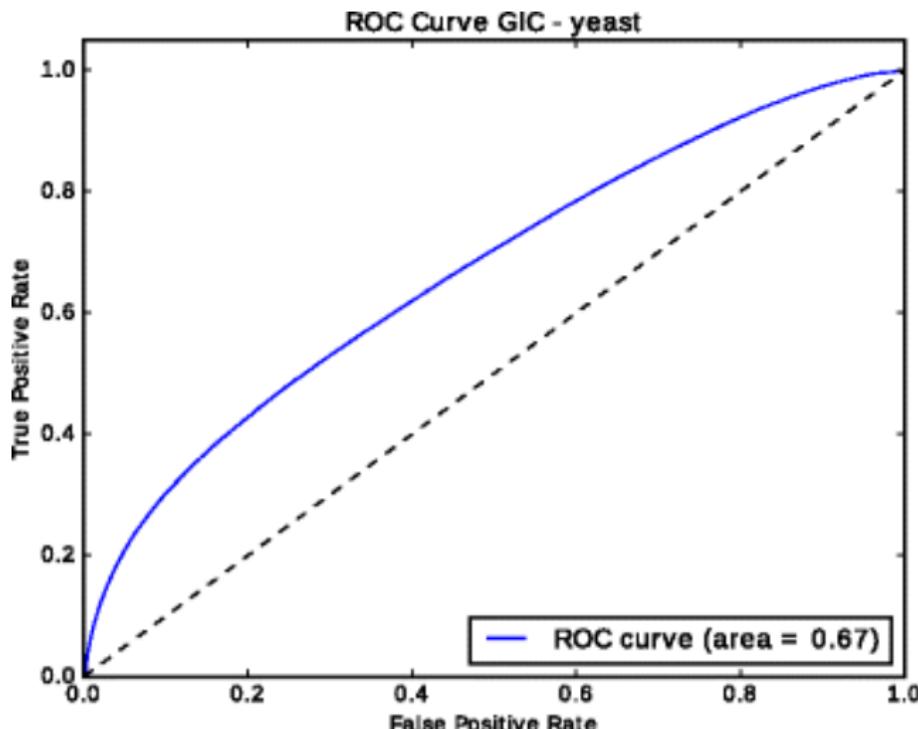
Applications of semantic similarity

Hypothesis

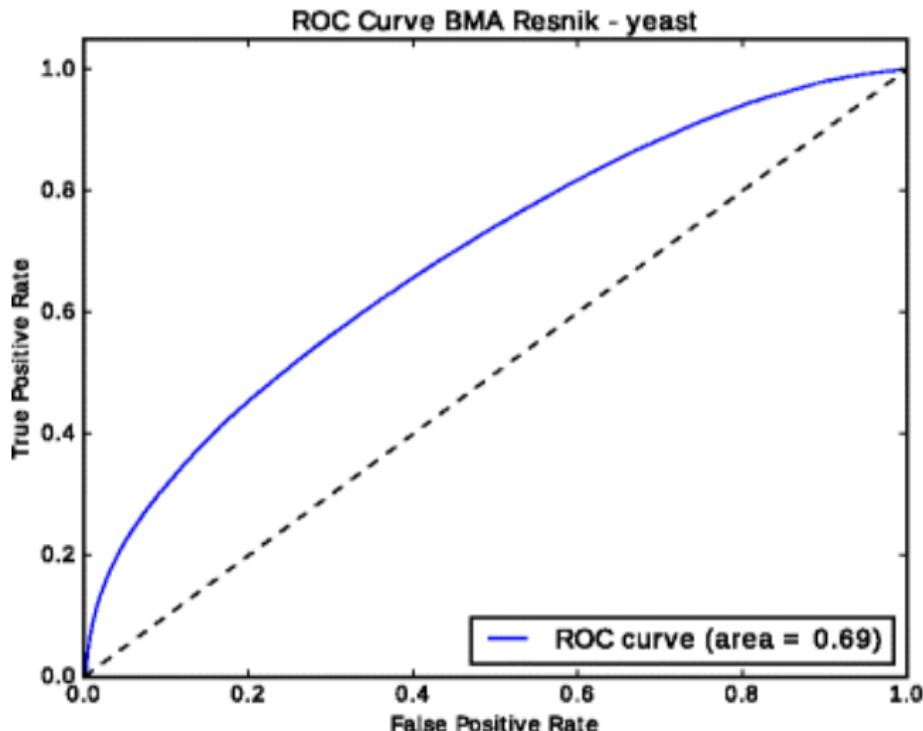
Interacting proteins have similar functions.

- ▶ relies on background knowledge about functions (encoded in GO)
- ▶ “similarity” can mean:
 - ▶ part of the same pathway
 - ▶ siblings of a common super-class
 - ▶ located in the same location
- ▶ set-based comparison of GO functions
 - ▶ single GO hierarchy or all?
 - ▶ which similarity measure?

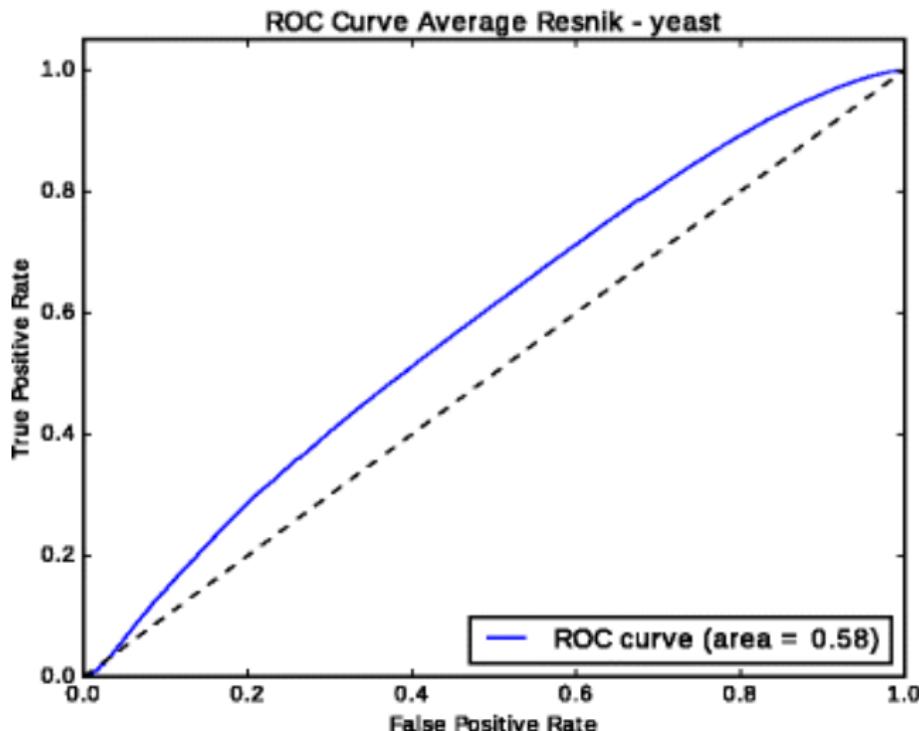
Applications of semantic similarity



Applications of semantic similarity



Applications of semantic similarity



Applications of semantic similarity

- ▶ no obvious choice of similarity measure
- ▶ depends on application
 - ▶ predicting PPIs in different organisms may benefit from a different similarity measure!
- ▶ different similarity measures may react differently to biases in data
- ▶ needs some testing and experience

Applications of semantic similarity

Recommendations:

- ▶ use Resnik's information content measure
- ▶ use Resnik's similarity
- ▶ use Best Match Average
- ▶ use the full ontology
- ▶ classify your ontology using an automated reasoner before applying semantic similarity
 - ▶ although many ontologies come pre-classified
- ▶ ⇒ but there are many exceptions
 - ▶ similar location ⇒ use location subset of GO
 - ▶ developmental phenotypes ⇒ use developmental branch of phenotype ontology

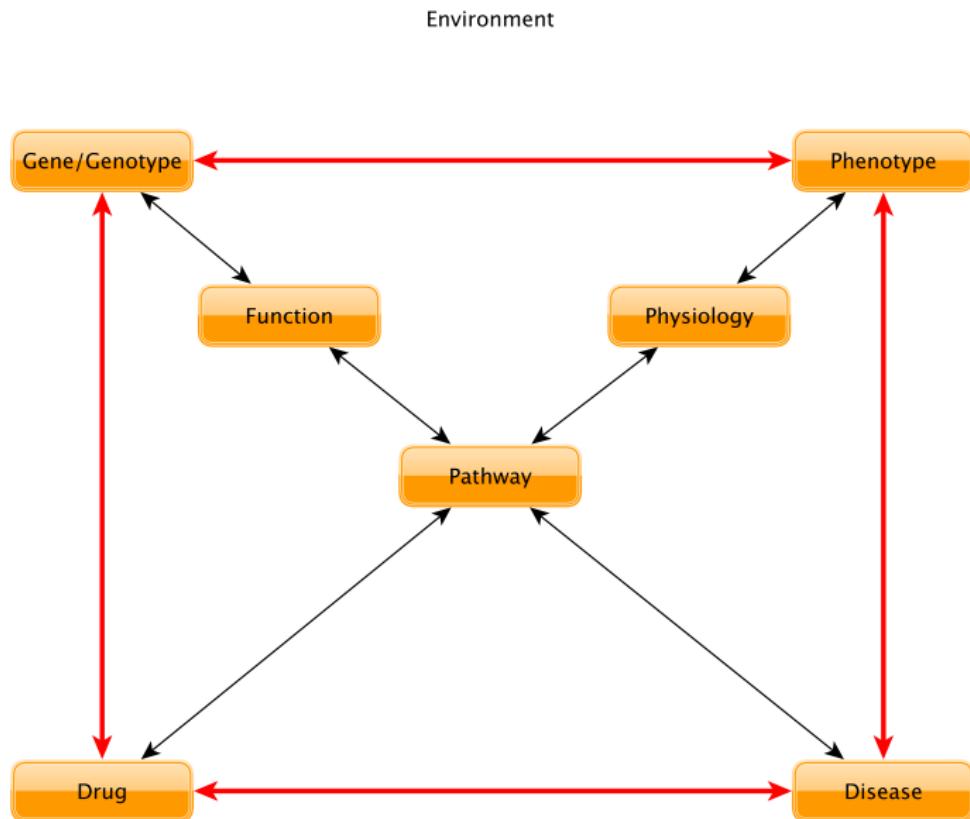
Applications of semantic similarity

- ▶ choice of ontology determines the kind of similarity
- ▶ functional similarity: Gene Ontology
- ▶ anatomical, structural similarity: anatomy ontologies (Uberon, MA, FMA, etc.)
- ▶ phenotypic similarity: phenotype ontology (HPO, MP, etc.)
- ▶ chemical structural similarity: ChEBI

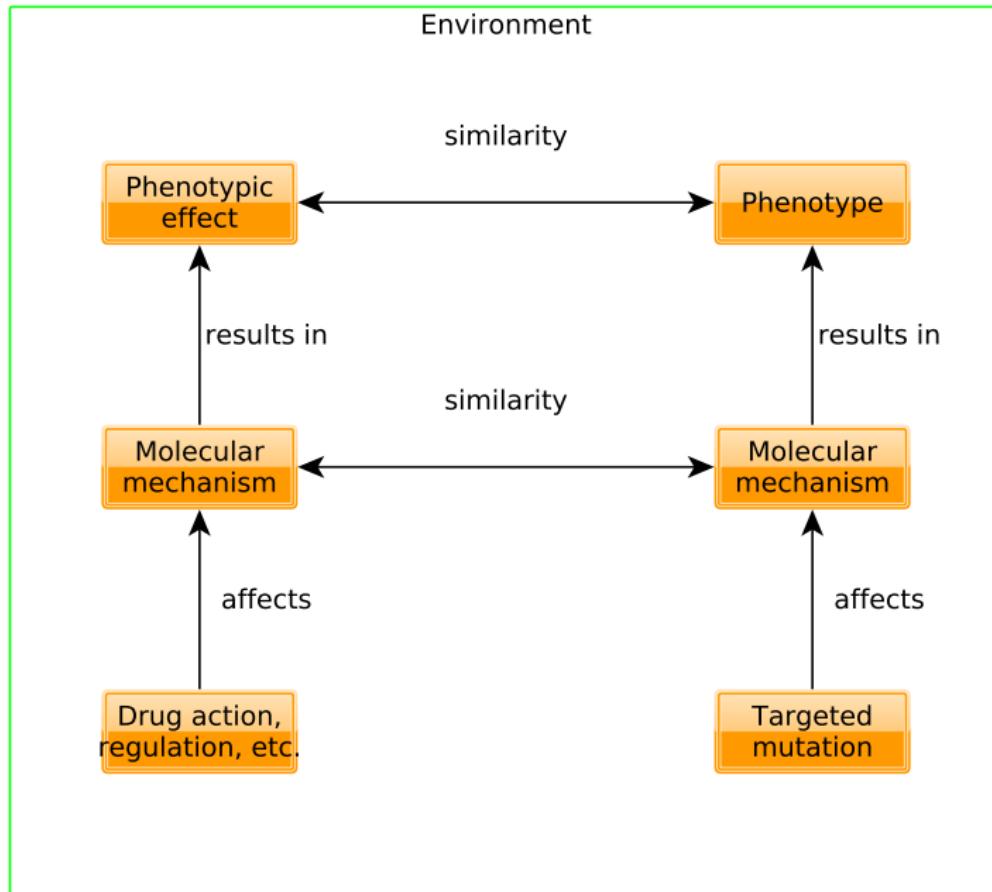
Applications of semantic similarity

- ▶ phenotypic similarity used to:
 - ▶ diagnosis: similarity between patient phenotypes and disease phenotypes
 - ▶ also between patient phenotypes and gene–phenotype associations
 - ▶ Phenomizer: <http://compbio.charite.de/phenomizer/>
 - ▶ disease modules: similarity between disease and disease
 - ▶ clustering/stratification: similarity between patient and patient
 - ▶ disease gene discovery: similarity between patient/disease phenotypes and gene–phenotype associations
 - ▶ in humans
 - ▶ in model organisms
 - ▶ drug repurposing: side-effect similarity; similarity between side effect profile and gene–disease associations

Applications of semantic similarity



Applications of semantic similarity



Applications of semantic similarity

- ▶ Guilt-by-association:
 - ▶ x is associated with y
 - ▶ z is similar to x
 - ▶ therefore: z may be associated with y
- ▶ candidate genes (polygenic disease):
 - ▶ FunSimMat: similar function \Rightarrow similar/same disease
 - ▶ side effect similarity: similar side effects \Rightarrow similar targets/indications

Applications of semantic similarity

- ▶ No guilt-by-association:
 - ▶ x causes a
 - ▶ y has b
 - ▶ a similar to b
 - ▶ therefore: b is caused by x
- ▶ candidate genes (monogenic and polygenic disease):
 - ▶ Phenomizer: gene x causes phenotypes a ; patient y has symptoms b ; a is similar to b ; therefore: gene x causes the symptoms in b
 - ▶ PhenomeNET: similar to Phenomizer but using model organism phenotypes (knockouts)
 - ▶ PhenomeDrug: knockout of gene x causes phenotypes a ; drug y causes side effects b ; a is similar to b ; therefore: drug y inhibits x (or: phenotypes b are caused by inhibition of x)
 - ▶ needs to compare model organism phenotypes and human phenotypes \Rightarrow ontology alignment/integration/mapping

Applications of semantic similarity

- ▶ comparing entities annotated with *different* ontologies/vocabularies of the *same* (or related) domains
 - ▶ medical: UMLS, HPO, DO, ORDO, NCIT, ICD, SNOMED CT, MeSH, ...
 - ▶ phenotype: HPO, MP, CPO, WBPhenotype, FBCV, MeSH, ...
 - ▶ chemical: ChEBI, MeSH, DrOn, RXNorm, DrugBank, ...
- ▶ needs mapping, alignment, or integration
 - ▶ mapping: given a term t , find corresponding class in ontology O
 - ▶ can be 1:1, 1:n, n:1, n:m
 - ▶ t can be from ontology, vocabulary, database, or text
 - ▶ use O for analysis
 - ▶ alignment: given two ontologies or vocabularies O_1 and O_2 , find all mappings between classes/terms in O_1 and O_2
 - ▶ applicable to ontologies and vocabularies
 - ▶ use O_1 or O_2 for analysis
 - ▶ integration: given two ontologies O_1 and O_2 , combine both ontologies into a single ontology O
 - ▶ maintain meaning of classes
 - ▶ use O for analysis

Applications of semantic similarity

- ▶ lexical mappings: use class labels (and synonyms) to find matches
 - ▶ hypertension (HP:0000822) and hypertension (MP:0000231)
- ▶ semantic mappings: use class axioms to find matches
 - ▶ pulmonary valve stenosis (MP:0006182) and Pulmonic stenosis (HP:0001642)
 - ▶ both definitions based on constricted (PATO:0001847) and pulmonary valve (UBERON:0002146)
- ▶ hybrid: combine lexical and semantic mappings

Applications of semantic similarity

tools for ontology mapping, matching, integration:

- ▶ AgreementMaker Light:
<https://github.com/AgreementMakerLight/AML-Jar>
 - ▶ structural (semantic) and lexical matches
 - ▶ can use domain-specific background knowledge
- ▶ LogMap: <https://github.com/ernestojimenezruiz/logmap-matcher>
 - ▶ structural (semantic) and lexical matches
 - ▶ biology-themed versions
- ▶ NCBO Annotator:
<https://bioportal.bioontology.org/annotator>
 - ▶ lexical matches only
 - ▶ can annotate full text
- ▶ recent tools and comprehensive ongoing evaluation:
 - ▶ OAEI: <http://oaei.ontologymatching.org/>

Applications of semantic similarity

semantic similarity and text mining:

- ▶ find all occurrences of classes of one (or more) ontologies in text
 - ▶ using lexical matching or semantic annotations of text
 - ▶ TextPresso (<http://www.textpresso.org/>), NCBO Annotator (<https://bioportal.bioontology.org/annotator>), WhatIzIt (<http://www.ebi.ac.uk/webservices/whatizit/info.jsf>)
 - ▶ ontology-specific text normalization tools
 - ▶ DNorm (diseases), GNorm (gene names), OSCAR (chemicals), ...
- ▶ use for database construction (automatic annotation), relation extraction, network construction (co-occurrence network), etc.

Applications of semantic similarity

<http://aber-owl.net/aber-owl/diseasephenotypes/>

- ▶ find phenotypes (signs and symptoms) associated with common diseases
 - ▶ no resource available for comparison
- ▶ pattern-based mining of literature with Aber-OWL: PubMed
- ▶ evaluation (of genetically based disease phenotypes) with experimentally validated disease genes

Applications of semantic similarity

<http://aber-owl.net/aber-owl/diseasephenotypes/>

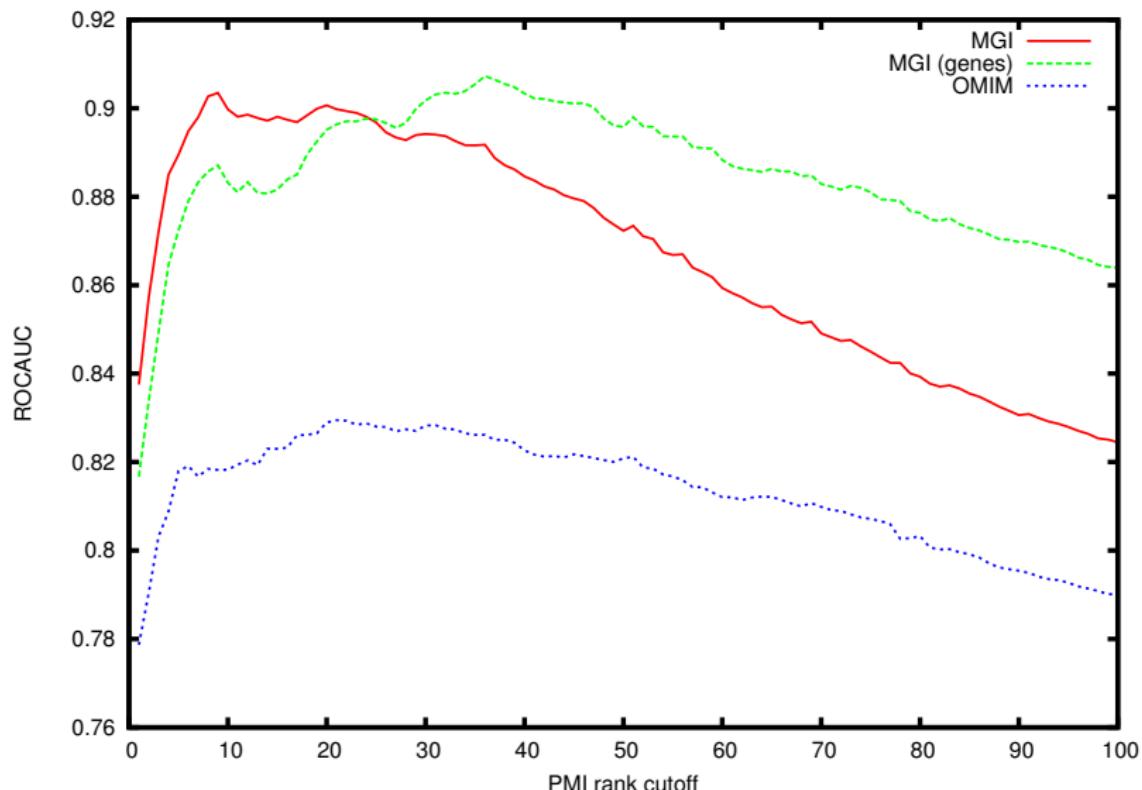
Phenotypes for 'bubonic plague'

- [Mouse models](#)
- [Network](#)

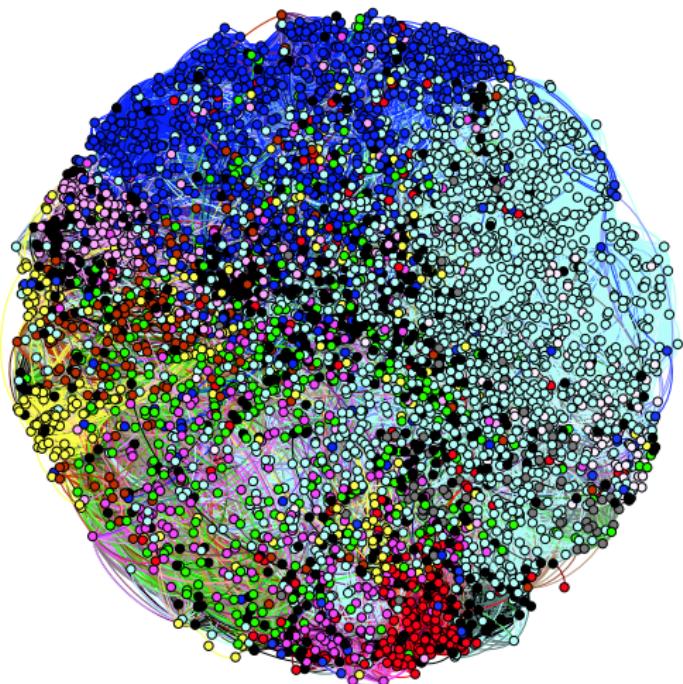
Show	100	entries	
	Phenotype	ID	Pointwise Mutual Information
death		MP:0001641	0.22837393990989352
Increased IgA level		HP:0003261	0.21745157416550895
eye lesions		MP:0013146	0.21110960811807034
Lymphadenitis		HP:0002840	0.20726444631165528
lymph node inflammation		MP:0003865	0.2070583004125834
Mediastinal lymphadenopathy		HP:0100721	0.16667331911036262
Abnormality of the lymph nodes		HP:0002733	0.14851361059698084
decreased susceptibility to infection		MP:0002409	0.1442019284611426
Pustule		HP:0200039	0.10968881998477503
tachypnea		MP:0005426	0.10850084511186676
Atrioventricular canal defect		HP:0006695	0.10364210532782134
Tachypnea		HP:0002789	0.10273022349597967
increased pulmonary respiratory rate		MP:0005573	0.09848893732226834

Applications of semantic similarity

<http://aber-owl.net/aber-owl/diseasephenotypes/>



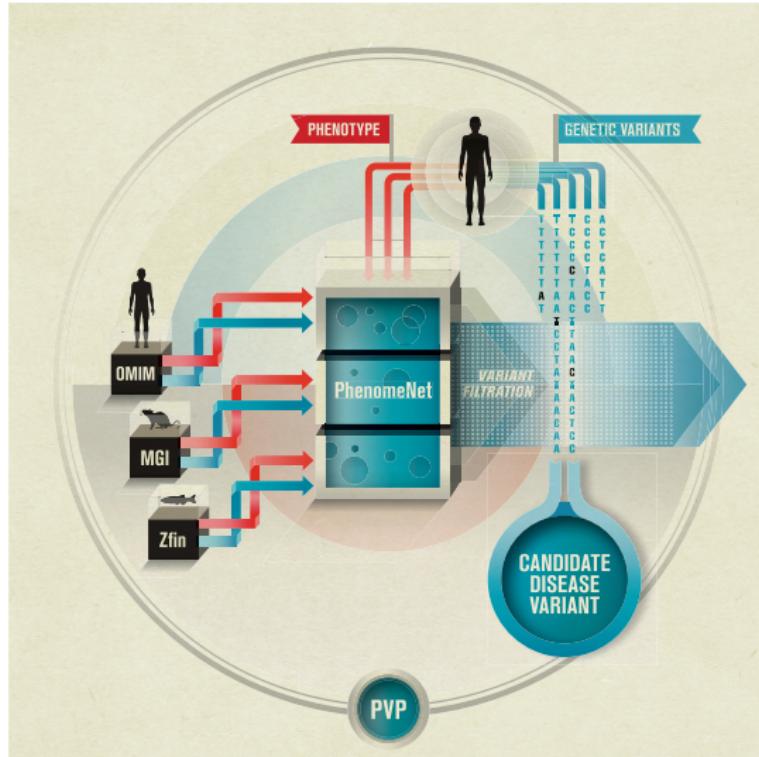
Applications of semantic similarity



Applications of semantic similarity

- ▶ semantic similarity can be used as features in machine learning models
 - ▶ when annotation space is too large
 - ▶ e.g., GO: 50,000 classes
 - ▶ replace binary representation
 - ▶ to incorporate background knowledge
 - ▶ semantic similarity encodes *implicitly* for ontology structure and axioms
 - ▶ encodes for *specificity* of classes
 - ▶ negative: reduce all annotations to single value
 - ▶ leads to loss of information
 - ▶ but is easier to use by many machine learning methods

Applications of semantic similarity



PVP Evaluation

- Synthetic data
 - 8,746 ClinVar pathogenic variants appended to WEX from 1000 Genome Project
 - 11,251 ClinVar pathogenic variants appended to WGS from 1000 Genome Project
- Real patient data
 - **19 exomes** from **UK10K_RARE_THYROID** study sample dataset.
 - 43 exomes from Addenbrooke's Hospital at Cambridge, UK with diseases of diverse phenotypes (e.g Bardett-Biedl Syndrome)
 - 1 patient's exome (Glutaric aciduria)
 - 5 genomes from Personal Genome Project

Synthetic WEX patients

	Top Hit	Top 10	ROC AUC
GWAVA	1.41%	6.32%	0.711
DANN	6.06%	26.69%	0.928
eXtasy	14.85%	42.99%	0.782
CADD	15.15%	32.05%	0.946
Exomiser	24.65%	58.56%	0.89
Phevor	28.25%	64.70%	0.912
PVP	78.80%	89.50%	0.978

Synthetic WGS patients

	Top Hit	Top 10	ROC AUC
GWAVA	0.46%	0.70%	0.949
DANN	4.95%	18.99%	0.99
CADD	5.76%	24.30%	0.962
Genomiser	31.35%	77.98%	0.948
PVP	76.47%	88.61%	0.995

Summary

- ▶ many semantic similarity measures
 - ▶ graph-based
 - ▶ feature-based
- ▶ useful for similarity-based prediction
 - ▶ similar entities ⇒ guilt-by-association
 - ▶ different entities
- ▶ combine with data and text mining
- ▶ features in machine learning methods

Acknowledgements

- ▶ Mona Alsharani
- ▶ Imene Boudellioua
- ▶ Maxat Kulmanov
- ▶ Sarah Alghamdi

Hands-on: semantic similarity

- ▶ if you have not done so *before* the tutorial, don't start now
 - ▶ you need to download *a lot* of data
 - ▶ you can just follow our demonstration and try later
 - ▶ (unless Internet is exceptionally fast for a conference Wifi, then just go ahead and do everything now)
- ▶ Jupyter Notebook
 - ▶ notebooks consist of code and rich text fragments
 - ▶ human readable (with nice figures) *and* executable
 - ▶ need to install the SciJava kernel (default: iPython)
 - ▶ very widely used
- ▶ Beaker Notebook
 - ▶ like Jupyter notebook, with lots of kernels pre-installed
 - ▶ Python, Groovy, Javascript, ...
- ▶ <https://github.com/bio-ontology-research-group/ontology-tutorial>

Hands-on: semantic similarity

In the tutorial, we will

- ▶ download an ontology
- ▶ explore the ontology with OWLAPI
- ▶ classify the ontology with an OWL reasoner
 - ▶ and query using an OWL reasoner
- ▶ store the inferred version locally
- ▶ use the Semantic Measures Library to:
 - ▶ explore the ontology as graph
 - ▶ compute similarity between classes
 - ▶ use different similarity measures
 - ▶ compare patients to mice
- ▶ you can build on this and extend for your own research!

Hands-on: semantic similarity

Do the tutorial...

Hands-on: semantic similarity

- ▶ now play with the Notebook:
 - ▶ look at the results list (check MGI)
 - ▶ try another disease (check OMIM)
 - ▶ or a drug effect (check SIDER)
- ▶ you can also test another ontology
 - ▶ GO for functional similarity
 - ▶ ChEBI for chemical (structural) similarity
 - ▶ or yeast phenotypes